

Geometry-Aware JEPA for EEG

A controlled frozen-transfer study on TUAB

Florent Tariolle Clément Genninasca Yoann Frayce Hippolyte du Pac de Marsoulies

Hack the World(s) – EEG Track · Team **HelloWorlds**

Contribution

Our contribution, in one slide

1. **A strong frozen baseline.** SIGReg-ambient JEPA matches the LeJEPA/Laya recipe's value on our setup – not a controlled head-to-head.
2. **Nothing beats it.** Ambient/tangent SIGReg & PEIRA, Graph-JEPA, Fourier-JEPA – none beats ambient SIGReg in balanced accuracy.
3. **Latent geometry is the contribution.** Visualised with **de Surrel's AIRM Riemannian t-SNE** (the right metric for SPD/EEG covariances).

Motivation

The question we actually asked

- **EEG foundation models** are everywhere – but most report *fine-tuned* numbers. What survives when you **freeze** the encoder?
- **JEPA** avoids collapse with an *anti-collapse regulariser*. The EEG covariance lives on a curved **SPD manifold**, not in flat Euclidean space.
- **Our intuition**: a *Riemannian geometry-aware* anti-collapse regulariser – acting in the SPD **tangent** of the EEG covariance, not flat ambient space – should *respect and improve* the latent geometry.
- **What we ship**: a strong frozen probe *and* a geometry-aware regulariser that **measurably sharpens the latent-space geometry**, visualised with de Surrél's AIRM Riemannian t-SNE.


Data

Data: TUH Abnormal EEG (TUAB)

- **Task:** binary *normal* vs *abnormal* EEG, recording level.
- **Signal:** 19 channels @ 200 Hz; 10 s windows (2,000 samples).
- **Standard split:** 2,717 train / 276 eval recordings, **patient-disjoint**.
- **Pretraining:** SSL reads only *unlabeled* train recordings.

Architecture

Two view and 1D convolutional encoder



What we tried


regulariser \ space	ambient (Euclidean)	tangent (SPD)
VICReg (variance/cov)	C0 reference	—
SIGReg (Gaussianity slices)	C1 <i>Laya-like baseline</i>	C2 geometry-aware
PEIRA (distribution-free)	C3	C4 <i>ex-hypothesis</i>

- **SIGReg** (LeJEPa/BCS): tests random 1-D slices for Gaussianity – assumes an *isotropic* target.
- **PEIRA**: distribution-free anti-collapse – no target distribution to mis-specify.
- **GraphJEPa**
- **FourierJEPa**

Results


SSL training curves (all cells healthy)

EEG-JEPA SSL training curves — 5 regularizer cells, 30 epochs, TUAB pretrain (seed 1)




Frozen & in-domain: we hold where frozen FMs fall

Frozen head-to-head on TUAB — FM rows: EEG-FM-Bench (arXiv 2508.17742, -



Beyond accuracy: frozen-probe robustness (3-seed)


Frozen-probe robustness on TUAB (3-seed mean $\pm$ std) — ambient SIGReg is the most dropout-robust



Visualization


Manifold-aware latent geometry (AIRM vs Euclidean)

Frozen SPD latents on TUAB eval — Euclidean (tangent) t-SNE vs AIRM Riemannian t-SNE
silhouette is full-dim (Euclidean tangent vs AIRM geodesic); 2-D maps are illustrative



Manifold viz generalises across sustained TUH tasks

Frozen SPD latents on TUSZ eval — 5 seizure types (cross-task, same-site)
Euclidean (tangent) t-SNE vs AIRM Riemannian t-SNE; silhouette is full-dim



Conclusion

Conclusion

- A **controlled frozen-transfer study**, not a leaderboard entry.
- A **strong frozen probe** (0.819 BA, ~ 0.89 AUROC) that holds where published FMs collapse when frozen.
- **The payoff is the latent geometry**: a Riemannian geometry-aware (SPD-tangent) regulariser **sharpens the frozen latent** – normal/abnormal and seizure classes form coherent structure – and **de Surrél's AIRM Riemannian embedding sharpens it further over the Euclidean view**, most clearly for SIGReg.
- On this simple binary task the probe accuracy is the same across variants – the tangent's gain is *geometric*, exactly where a manifold-aware regulariser should help.

References

References i

- Balestrieri, R. & LeCun, Y. (2025). *LeJEPA: Provable and Scalable Self-Supervised Learning Without the Heuristics*. Preprint. (*SIGReg – our anti-collapse objective*)
- Panchavati, S. et al. (2026). *Laya: a LeJEPA approach to EEG via latent prediction over reconstruction*. Preprint. (*our ambient baseline*)
- Arbel, M., Terver, B. & Ponce, J. (2026). *PEIRA: Learning Predictive Encoders through Inter-View Regressor Alignment*. Preprint.
- Bardes, A., Ponce, J. & LeCun, Y. (2022). *VICReg*. ICLR.
- LeCun, Y. (2022). *A Path Towards Autonomous Machine Intelligence (JEPA)*. OpenReview.
- de Surrél, T., Lotte, F., Chevallier, S. & Yger, F. (2025). *Wrapped Gaussian on the manifold of Symmetric Positive Definite Matrices*. Preprint.

References ii

- de Surrél, T. *r_tSNE_in_C* – Riemannian t-SNE on SPD matrices (AIRM). github.com/thibaultdesurrel/r_tSNE_in_C. (*our manifold-aware latent viz*)
- Wang, Z. et al. (2025). *EEG-ReMinD: Self-Supervised State Reconstruction-Primed Riemannian Dynamics*. Preprint. (*Riemannian SPD SSL; reconstruction, not JEPA*)
- Chen, M. et al. (2025). *MENDR: Manifold Explainable Neural Data Representations*. Preprint. (*first Riemannian-SPD EEG FM; precedent for the geometry null*)
- Gemein, L. et al. (2020). *Machine-learning-based diagnostics of EEG pathology*. NeuroImage. (*Riemannian baseline*)
- Jiang, W.-B. et al. (2024). *LaBraM: Large Brain Model for EEG*. ICLR.
- Wang, J. et al. (2024). *CBraMod: Criss-Cross Brain Foundation Model*. Preprint.
- Yang, C. et al. (2023). *BIOT: Biosignal Transformer*. NeurIPS.

- Wang, G. et al. (2024). *EEGPT: Pretrained Transformer for Universal and Reliable Representation of EEG Signals*. NeurIPS.
- Foumani, N. et al. (2024). *EEG2Rep: Enhancing Self-supervised EEG Representation Through Informative Masked Inputs*. KDD.
- Kostas, D., Aroca-Ouellette, S. & Rudzicz, F. (2021). *BENDR: Transformers & Contrastive SSL for EEG*. Front. Hum. Neurosci.
- Xiong, W. et al. (2026). *EEG-FM-Bench: Systematic Evaluation & Diagnostic Analyses of EEG Foundation Models*. ICML. (source of our frozen FM head-to-head numbers)
- Obeid, I. & Picone, J. (2016). *The Temple University Hospital EEG Data Corpus*. Front. Neurosci. 10:196. (TUAB / TUSZ / TUEV)

- eb_jepa – JEPA reference implementation, Meta AI / FAIR.
github.com/facebookresearch/eb_jepa. (*our trimmed code vendor*)