Geometry-Aware JEPA for EEG

A controlled frozen-transfer study on TUAB

Florent Tariolle Clément Genninasca Yoann Frayce Hippolyte du Pac de Marsoulies

Hack the World(s) – EEG Track · Team **HelloWorlds**

Contribution

Our contribution, in one slide

1. **A strong frozen baseline.** SIGReg-ambient JEPA matches the LeJEPA/Laya recipe's value on our setup – not a controlled head-to-head.
2. **Nothing beats it.** Ambient/tangent SIGReg & PEIRA, Graph-JEPA, Fourier-JEPA – none beats ambient SIGReg in balanced accuracy.
3. **Latent geometry is the contribution.** Visualised with **de Surrel's AIRM Riemannian t-SNE** (the right metric for SPD/EEG covariances).

Motivation

The question we actually asked

- **EEG foundation models** are everywhere – but most report *fine-tuned* numbers. What survives when you **freeze** the encoder?
- **JEPA** avoids collapse with an *anti-collapse regulariser*. The EEG covariance lives on a curved **SPD manifold**, not in flat Euclidean space.
- **Our pre-registered hypothesis**: making anti-collapse *geometry-aware* (act in the SPD **tangent**) and *distribution-free* (**PEIRA**) beats a plain ambient baseline.
- **What we ship**: a controlled study with an *honest* answer – a strong frozen probe, a clean negative result with a geometric mechanism, and latent space visualizations.

Data

Data: TUH Abnormal EEG (TUAB)

- **Task**: binary *normal* vs *abnormal* EEG, recording level.
- **Signal**: 19 channels @ 200 Hz; 10 s windows (2,000 samples).
- **Standard split**: 2,717 train / 276 eval recordings, **patient-disjoint**.
- **Pretraining**: SSL reads only *unlabeled* train recordings.

Architecture

Two view and 1D convolutional encoder

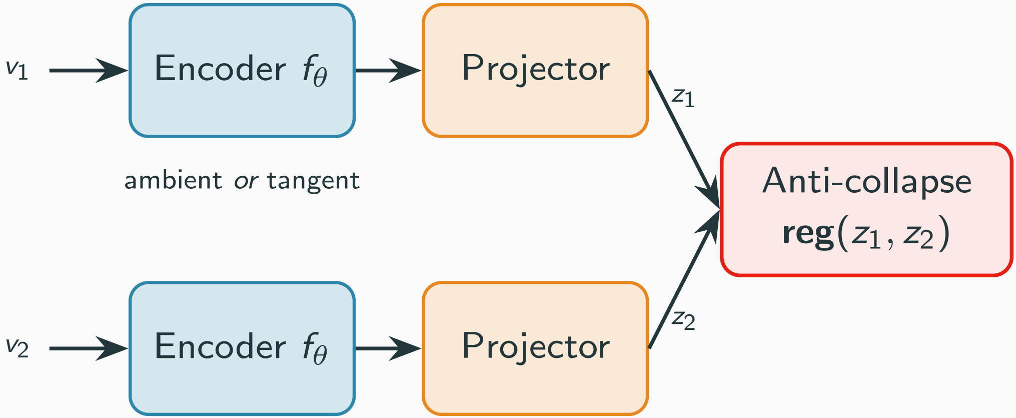

What we tried

regulariser \ **space**	**ambient** (Euclidean)	**tangent** (SPD)
VICReg (variance/cov)	`C0` reference	—
SIGReg (Gaussianity slices)	`C1` *Laya-like baseline*	`C2` geometry-aware
PEIRA (distribution-free)	`C3`	`C4` *ex-hypothesis*

- **SIGReg** (LeJEPA/BCS): tests random 1-D slices for Gaussianity – assumes an *isotropic* target.
- **PEIRA**: distribution-free anti-collapse – no target distribution to mis-specify.
- **GraphJEPA**
- **FourierJEPA**

Results

SSL training curves (all cells healthy)

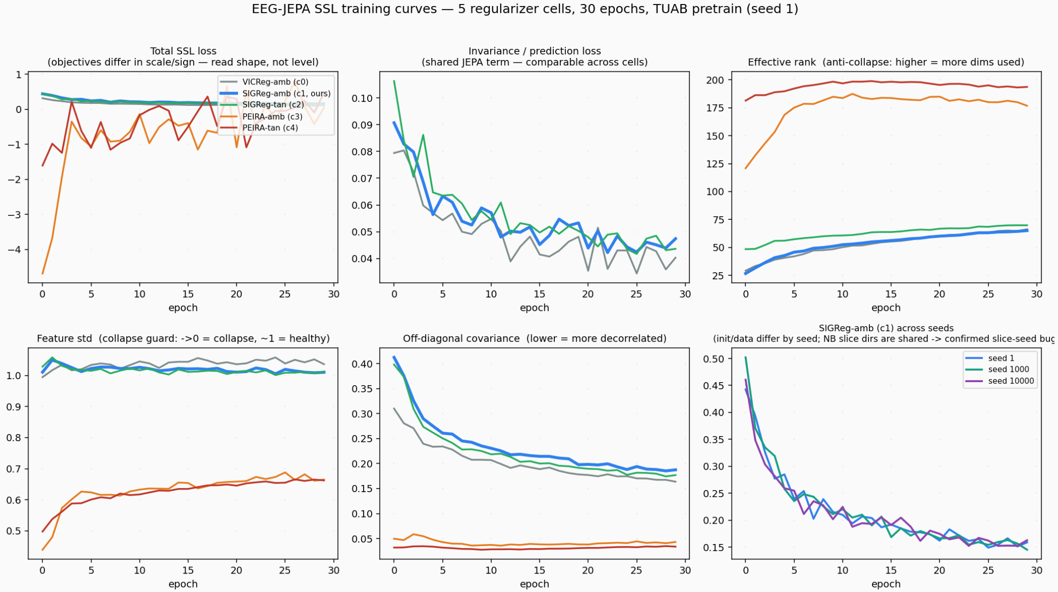

Frozen & in-domain: we hold where frozen FMs fall

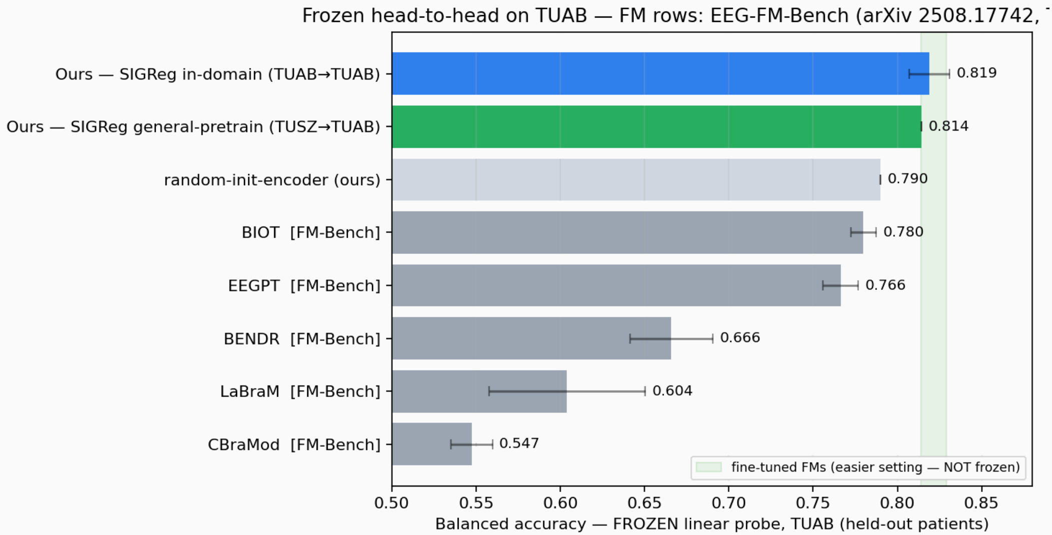

Beyond accuracy: frozen-probe robustness (3-seed)

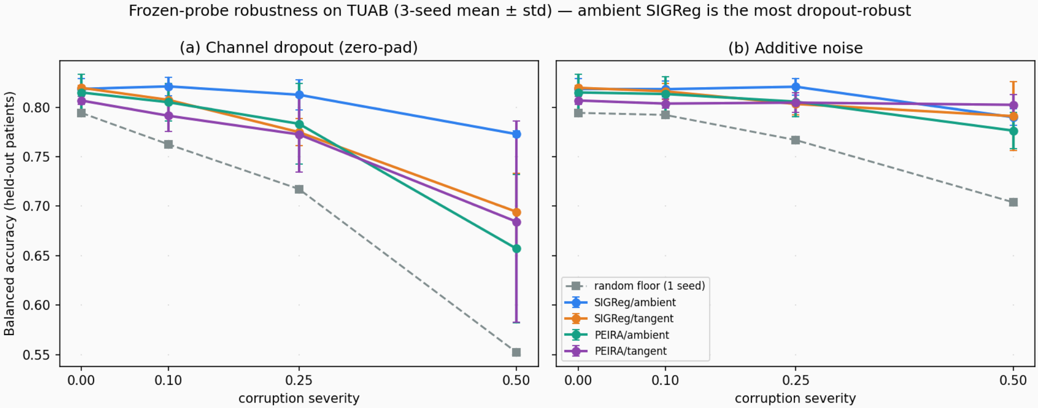

Visualization

Manifold-aware latent geometry (AIRM vs Euclidean)

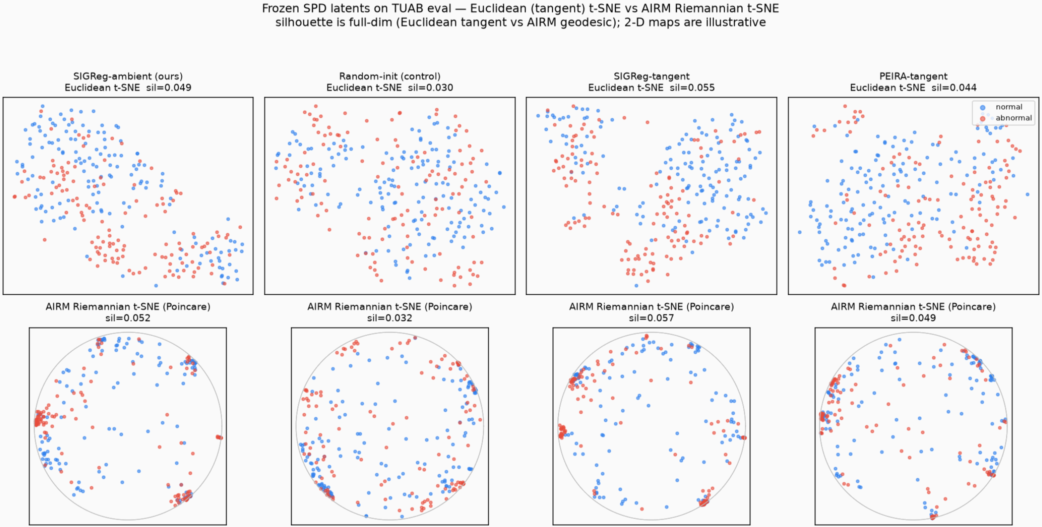

Manifold viz generalises across sustained TUH tasks

Conclusion

Conclusion

- A **controlled frozen-transfer study**, not a leaderboard entry.
- **A strong frozen probe** (0.819 BA, $\sim$0.89 AUROC) that holds where published FMs collapse when frozen.
- **The null is multi-axis**: geometry/PEIRA help neither accuracy, nor calibration, nor robustness (3 seeds).
- **The payoff is the latent geometry**: the frozen SPD latent is visibly *organised* – normal/abnormal and seizure classes form coherent structure – and **de Surrel's AIRM Riemannian embedding sharpens it over the Euclidean view on every clustered task, most clearly for SIGReg**.

Built on a trimmed `eb_jepa` vendor. Team **Hello Worlds**.

References

References i

- Balestriero, R. & LeCun, Y. (2025). *LeJEPA: Provable and Scalable Self-Supervised Learning Without the Heuristics*. Preprint. *(SIGReg – our anti-collapse objective)*
- Panchavati, S. et al. (2026). *Laya: a LeJEPA approach to EEG via latent prediction over reconstruction*. Preprint. *(our ambient baseline)*
- Arbel, M., Terver, B. & Ponce, J. (2026). *PEIRA: Learning Predictive Encoders through Inter-View Regressor Alignment*. Preprint.
- Bardes, A., Ponce, J. & LeCun, Y. (2022). *VICReg*. ICLR.
- LeCun, Y. (2022). *A Path Towards Autonomous Machine Intelligence (JEPA)*. OpenReview.
- de Surrel, T., Lotte, F., Chevallier, S. & Yger, F. (2025). *Wrapped Gaussian on the manifold of Symmetric Positive Definite Matrices*. Preprint.

References ii

- de Surrel, T. *r_tSNE_in_C* – Riemannian t-SNE on SPD matrices (AIRM). github.com/thibaultdesurrel/r_tSNE_in_C. *(our manifold-aware latent viz)*
- Wang, Z. et al. (2025). *EEG-ReMinD: Self-Supervised State Reconstruction-Primed Riemannian Dynamics*. Preprint. *(Riemannian SPD SSL; reconstruction, not JEPA)*
- Chen, M. et al. (2025). *MENDR: Manifold Explainable Neural Data Representations*. Preprint. *(first Riemannian-SPD EEG FM; precedent for the geometry null)*
- Gemein, L. et al. (2020). *Machine-learning-based diagnostics of EEG pathology*. NeuroImage. *(Riemannian baseline)*
- Jiang, W.-B. et al. (2024). *LaBraM: Large Brain Model for EEG*. ICLR.
- Wang, J. et al. (2024). *CBraMod: Criss-Cross Brain Foundation Model*. Preprint.
- Yang, C. et al. (2023). *BIOT: Biosignal Transformer*. NeurIPS.

References iii

- Wang, G. et al. (2024). *EEGPT: Pretrained Transformer for Universal and Reliable Representation of EEG Signals*. NeurIPS.
- Foumani, N. et al. (2024). *EEG2Rep: Enhancing Self-supervised EEG Representation Through Informative Masked Inputs*. KDD.
- Kostas, D., Aroca-Ouellette, S. & Rudzicz, F. (2021). *BENDR: Transformers & Contrastive SSL for EEG*. Front. Hum. Neurosci.
- Xiong, W. et al. (2026). *EEG-FM-Bench: Systematic Evaluation & Diagnostic Analyses of EEG Foundation Models*. ICML. *(source of our frozen FM head-to-head numbers)*
- Obeid, I. & Picone, J. (2016). *The Temple University Hospital EEG Data Corpus*. Front. Neurosci. 10:196. *(TUAB / TUSZ / TUEV)*

References iv

- `eb_jepa` – JEPA reference implementation, Meta AI / FAIR. github.com/facebookresearch/eb_jepa. *(our trimmed code vendor)*